

CS&SS 321 - Data Science and Statistics for Social Sciences

Module I - Getting started with R/RStudio

Ramses Llobet

Welcome!



(But Seriously) Welcome!

- ▶ Welcome to the first quiz section of CS&SS / SOC / STAT 321!
- ▶ I am Ramses Llobet (rllobet@uw.edu), I am a Ph.D. candidate in Political Science.
- ▶ My research interest are in **political economy** and **applied statistics**.
- ▶ Please **DO NOT** hesitate to stop me if you don't hear or understand me properly.
- ▶ **DO NOT** hesitate to ask questions. No question is silly. :)

Now it's your turn

- ▶ **Name** and major/year (or intended major)?
- ▶ **Why** are you take this course?
- ▶ **What** is your experience with R (zero shame)?

R setup

- ▶ How to install R and R-studio.
 - ▶ R-4.3.2 for [Windows](#)
 - ▶ R-4.3.2 for [macOS](#)
- ▶ R-studio can be downloaded from [posit's repository](#).
- ▶ I recommend this [tutorial](#) from Casey Bates for an overview of R and RStudio.
- ▶ **Live coding**: how to install packages and start tutorials (`setting_up.R`).
- ▶ For Mac users, installation of the **qss** package may sometimes fail if **pandoc** or **curl** is not installed or updated on your Mac. To resolve this, you can:
 1. Install the package manager [Homebrew](#) package.
 2. Then use the macOS terminal to install **pandoc** or **curl** using the commands `brew install pandoc` or `brew install curl`.

Useful free online R resources

- ▶ Introductory:
 - ▶ Grolemund ([2014](#)) *Hands-On Programming with R*.
- ▶ Intermediary:
 - ▶ Wickham et al. ([2023](#)) *R for Data Science*. 2nd Edition.
- ▶ R Markdown
 - ▶ Xie et al. ([2022](#)) *R Markdown: The Definitive Guide*
- ▶ Others
 - ▶ [Stack Overflow](#).
 - ▶ [ChatGPT](#)

Project management and working directory

- ▶ A good practice is to keep your projects and files organized and tidy.
 - ▶ **Avoid** accumulating data and R files in your **downloads folder**.
- ▶ I recommend creating an **R project** file in your course folder materials. R projects have several advantages:
 - ▶ Centralized and efficient *workflow*.
 - ▶ Sets the **current** (*root*) working directory.
 - ▶ See more in Martin Chan's [beginner's guide](#).

What are working directories?

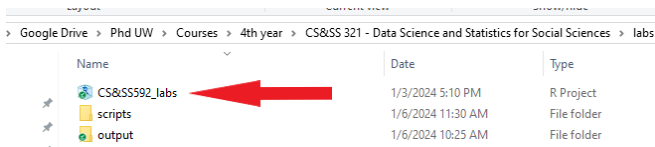
- ▶ A **directory** is a **folder** in a file system that stores files and other sub-directories.
- ▶ A **path** is a string that specifies the **location** of a directory in a file system.
- ▶ For example:
 - ▶ D:\Google Drive
 - ▶ D:\Google Drive\Phd UW\Courses\Third Year\CSSS 594
- Text as Data
- ▶ When you **run** a command or script, R looks for files and sub-directories based on **relative paths** to your current working directory.

Absolute and relative paths

- ▶ **Absolute Path:** Specifies the full path from the **root** directory to the file or directory.
 - ▶ For example: D:\Google Drive\Phd UW\Courses\Third Year\CSSS 321\scripts\setting_up.R is an **absolute path**.
- ▶ **Relative Path:** Specifies the path relative to the current working directory.
 - ▶ For example, if the **working directory** is D:\Google Drive\Phd UW\Courses\Third Year\CSSS 321, then
 - ▶ scripts\setting_up.R is a **relative path**.

Project management: working directory

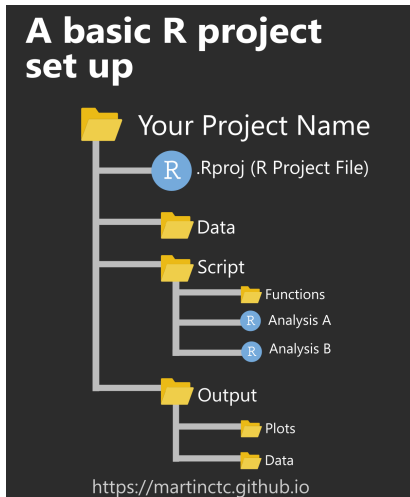
- ▶ **.Rproj** (R Project File) in your project folder establishes the working directory as its absolute path upon opening R.



Google Drive > Phd UW > Courses > 4th year > CS&SS 321 - Data Science and Statistics for Social Sciences > labs		
Name	Date	Type
CS&SS592_labs	1/3/2024 5:10 PM	R Project
scripts	1/6/2024 11:30 AM	File folder
output	1/6/2024 10:25 AM	File folder

- ▶ Employing **.Rproj** and **relative paths** in R streamlines project management and collaboration by overseeing files, inputs, and outputs.
 - ▶ **Live demonstration** of how to create and manage an R Project File.

Project management: workflow



Working directories: obsolete practices

- ▶ Workflow with *.Rproj* is relatively **new**.
- ▶ Until recently, users had to **manually set** working directories using **functions** or specialized **packages**. See example:

```
setwd() # function to set directory
```

```
setwd("D:/Google Drive/Phd UW/Courses  
/Third Year/CSSS 594 - Text as Data  
/presentation") # remember to put quotes
```

What are functions?

- ▶ They are a **set of instructions** that performs a specific task in R.
- ▶ Functions often take one or more **arguments**, which are inputs that are used to customize the behavior of the function.
- ▶ The `mean()` function takes one **required** argument, which is the vector of numbers to calculate the mean of.

```
# create a vector consisting of midterm scores.  
grades_M <- c(76, 82, 94, 45, 75)
```

```
# calculate the mean using the mean() function  
mean(grades_M)
```

```
## [1] 74.4
```

What are functions?

- ▶ the `mean()` function also has additional optional arguments, which can be used to further customize the behavior of the function.

```
# create a vector consisting of final scores.  
grades_F <- c(82, 90, 89, NA, 64)  
  
# calculate the mean using the mean() function  
mean(grades_F)
```

```
## [1] NA
```

```
# use the argument `na.rm` to evaluate the removal of NAs  
mean(grades_F, na.rm= TRUE)
```

```
## [1] 81.25
```

- ▶ **Remember:** use `?` or `help()` to see the documentation of a function.

R-Markdown

- ▶ Save the following [Cheat Sheet](#) for RMarkdown.
- ▶ If any of you is looking for an general introduction for RMarkdown, I suggest you to check [Chapter 27](#) from Wickham and Grolemund ([2017](#)) - **R for Data Science**.
- ▶ If you want a more comprehensive guide, then check Xie et al. ([2021](#)) - **R Markdown: The Definitive Guide**.
- ▶ Another, more applied, resource is Xie et al. ([2022](#)) - **R Markdown Cookbook**.

R-Markdown

- ▶ RMarkdown is a document format that allows you to integrate R **code** and **output** into a single document.
- ▶ Besides R code and output, it can also include **text**, **images**, and other **multimedia elements**, allowing for rich and informative documents.
- ▶ *Pandoc* is a free and open-source **document converter** that can convert documents from one markup language to another.
 - ▶ In the context of Rmarkdown, pandoc is the underlying document converter (software) that converts the R-markdown file into a final output format, such as **HTML**, **PDF**, or **Word**.

R-Markdown

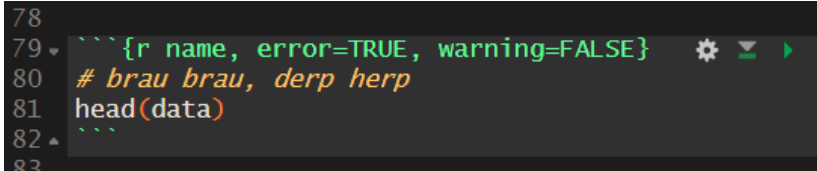
- The output format of the final document can be customized using options in the **YAML header** or external templates.

```
1 ---
2 title: "Lab 1 - Intro to RMarkdown"
3 author: "Your name"
4 date: \today
5 output:
6   pdf_document:
7     latex_engine: pdflatex
8   fontsize: 12pt
9   editor_options:
10     chunk_output_type: console
11 ---
12
```

- The YAML header in RMarkdown is a block of configuration settings at the beginning of the document enclosed by three hyphens (---).
- It is used to specify document metadata and other settings such as the document title, author, output format, and more.

R-Markdown

- ▶ **Code chunks** are sections of R code that can be executed and embedded within an RMarkdown document.



```
78  
79 ```{r name, error=TRUE, warning=FALSE}  
80 # brau brau, derp herp  
81 head(data)  
82 ```  
83
```

- ▶ Code chunks can be inserted using the syntax `{r}` and closed with ````.
 - ▶ Short cut in Windows: Ctrl + Alt + I
 - ▶ Short cut in macOS: Cmd + Option + I`
- ▶ Code chunks can be customized with various **chunk options**.
- ▶ **Note:** set the function `knitr::opts_chunk$set()` with any general setting without repeating it in every code chunk.

R-Markdown


► Frequently used chunk options

Option	Description
include	If FALSE, knitr will run the chunk but not include the chunk in the final document
echo	If FALSE, knitr will not display the code in the code chunk above it's results in the final document.
error	If FALSE, knitr will not display any error messages generated by the code.
message	If FALSE, knitr will not display any messages generated by the code.
warning	If FALSE, knitr will not display any warning messages generated by the code.

Recommendation for Homework

Option	HW setting
include	TRUE
echo	TRUE
error	FALSE
message	FALSE
warning	FALSE

```
1 ---
2 title: "RMarkdown sample"
3 author: "Your name"
4 date: "2024-01-10"
5 output: pdf_document
6 ---
7
8 ```{r setup, include=FALSE}
9 # This first chunk is generally hidden and used to load data, libraries
10 # and the stuff that you do not need to show in the report.
11 knitr::opts_chunk$set(echo = TRUE,
12                       error = FALSE,
13                       message = FALSE,
14                       warning = FALSE)
15
16 # Load libraries
17
18 library("tidyverse")
19
```

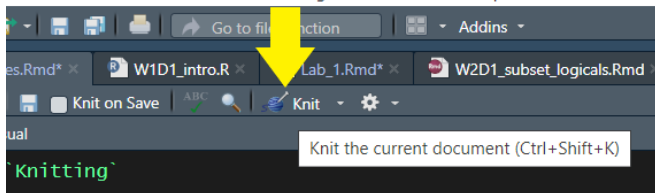


R-Markdown

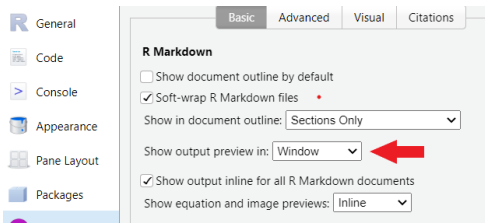
- ▶ In RMarkdown, **rendering** a document means converting the source RMarkdown file into its final output format (using pandoc).
- ▶ To render a document, we need to Knit, knitting is the process of taking the RMarkdown file and converting it into a single, cohesive document that can be rendered into different formats (HTML, PDF, etc).
 - ▶ To produce **PDF file**, you need TeX files.
- ▶ Easy way: Install the tinytex package:
`install.packages("tinytex")`. Then run
`tinytex::install_tinytex()`.

Knitting

- To knit:

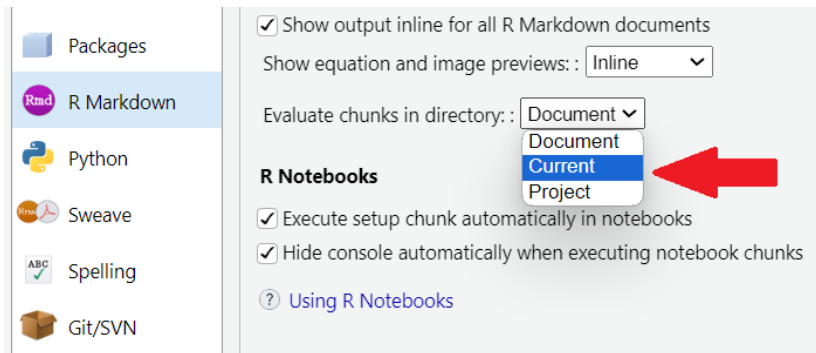


- Auxiliary window for output preview:



Working directories and R-Markdown

- ▶ When opening an RMarkdown file, this will set the file location as the **working directory**.
- ▶ Change the following option in the **global options** to avoid this behavior:



The screenshot shows the 'Global Options' dialog box in RStudio. On the left, a sidebar lists various options: Packages, R Markdown (selected), Python, Sweave, Spelling, and Git/SVN. The main panel displays settings for 'R Markdown' documents. The 'Show output inline for all R Markdown documents' checkbox is checked. Below it, 'Show equation and image previews:' is set to 'Inline'. The 'Evaluate chunks in directory:' dropdown menu is open, showing three options: 'Document', 'Current' (highlighted with a blue bar and a red arrow pointing to it), and 'Project'. Under the 'R Notebooks' section, two checkboxes are checked: 'Execute setup chunk automatically in notebooks' and 'Hide console automatically when executing notebook chunks'. At the bottom, there is a link with a question mark icon labeled 'Using R Notebooks'.

Packages

R Markdown

Python

Sweave

Spelling

Git/SVN

☒ Show output inline for all R Markdown documents

Show equation and image previews: : Inline

Evaluate chunks in directory: : Document

Document

Current

Project

R Notebooks

☒ Execute setup chunk automatically in notebooks

☒ Hide console automatically when executing notebook chunks

? Using R Notebooks

R-Markdown

- ▶ Live demonstration and in-class exercise:
 - ▶ Open the file `RMarkdown_sample.Rmd`

Getting help: reading the error messages

Reference: [David Robinson](#)

```
x <- 10 + foo
```

Error: object 'foo' not found: You tried to access a variable that doesn't exist.

You might have:

- ▶ **misspelled** the variable name
- ▶ incorrectly **capitalized** the variable name (R is case sensitive!)
- ▶ **forgotten** to run the line that defines the variable in the first place, or run into an error on that line.

Getting help: reading the error messages

```
x <- foo(...)
```

Error: could not find function "foo": You tried to use a function that doesn't exist. You might have:

- ▶ **misspelled** the function name
- ▶ incorrectly **capitalized** the function name
- ▶ forgotten to **load the library** that provides this function.

Getting help: reading the error messages

```
x <- c(1:10))
```

Error: unexpected ')' in ...: There is an extra end parenthesis in your line

```
x <- 10; y <- 20  
mean(x y)
```

Error: unexpected symbol in ...: The most common cause of this is forgetting a punctuation mark such as a comma: for example, `foo(bar1 bar2)` instead of `foo(bar1, bar2)`.

Error: unexpected numeric constant is similar: it just means the value after the missing punctuation is a number (for example, `x 2` instead of `x = 2`).

Getting help: reading the error messages

```
paste("welcome to CSSS, 321)
```

- ▶ You might see a + sign in the interpreter after you hit return. This means the previous statement is unfinished:
- ▶ it might have an **open parenthesis** that never closes,
 - ▶ an open " or ' that is unmatched, or
 - ▶ it could end with an operator like + or - that expects the line to continue afterwards.
- ▶ Find the problem in your previous lines (count parentheses, and check your quotes) and fix it.

Getting help: reproducible code

- ▶ If you feel stuck with an error, seek help but remember to provide **reproducible code** in an R-script file:
 1. Load necessary **packages** at the beginning.
 2. Include all code up to the error, or at least the **necessary** to reproduce it.
 3. **Comment** your code for clarity.
 4. If applicable, send the necessary **data** to reproduce the error.

Review of basics

- ▶ The subsequent slides show some R basics that we have already covered during the live demonstrations but feel free to review them on your own.

Running R code and operators

```
# Arithmetic Operators
```

```
1 + 1
```

```
## [1] 2
```

```
2 * 8
```

```
## [1] 16
```

```
9 / 3
```

```
## [1] 3
```

```
2^3
```

```
## [1] 8
```

Running R code and operators

```
# Relational Operators
```

```
10 > 8 # is 10 bigger than 8?
```

```
## [1] TRUE
```

```
7 <= 6 # is 7 less or equal to 6?
```

```
## [1] FALSE
```

```
(2 * 5) == 10 # is 2*5 equal to 10?
```

```
## [1] TRUE
```

```
1 != 2 # is 1 unequal to 2?
```

```
## [1] TRUE
```


Objects in R: vectors and assignment

```
# Concatenate vectors into a new vector  
c(1, 2, 3)
```

```
## [1] 1 2 3
```

```
# Assign them to a new object for manipulation  
x <- c(1, 2, 3)  
print(x) # or simply, x
```

```
## [1] 1 2 3
```

```
# Operators on vector  
x + 1
```

```
## [1] 2 3 4
```

```
# Logical test on vector  
x == 1
```

Objects in R: vectors and functions

Use an object as input to a function

```
x <- c(1, 2, 3)
```

Functions take input(s) and produce output(s)

```
class(x)
```

```
## [1] "numeric"
```

```
length(x)
```

```
## [1] 3
```

```
mean(x)
```

```
## [1] 2
```

Objects in R: introductory tips

- Unless you assign (`<-`) some operations or transformations to an object, those values will not be registered

```
x <- c(1, 2, 3)
print(x + 1)
```

```
## [1] 2 3 4
```

```
print(x)
```

```
## [1] 1 2 3
```

```
x <- x + 1
print(x)
```

```
## [1] 2 3 4
```

Objects in R: introductory tips

- New assignment will overwrite the original values if you assign some values to an existing object. It is a **major** source of errors. One advise is to keep distinct object names

```
x <- c(1, 2, 3)
length(x)
```

```
## [1] 3
```

```
x <- c(1, 2, 3, 4, 5)
length(x)
```

```
## [1] 5
```

Objects in R: atomic vectors

- Most common types of atomic vectors: **numeric (integer, double)**, **logical**, **character**

```
x <- c(1, 2, 3)
class(x)
```

```
## [1] "numeric"
```

```
y <- c(TRUE, FALSE, FALSE)
class(y)
```

```
## [1] "logical"
```

```
names <- c("Peter", "Paul", "Mary")
class(names)
```

```
## [1] "character"
```

Objects in R: atomic vectors

- You can also coerce one type of vector into another:

```
x <- c(1, 2, 3)
x <- as.character(x)

print(x)
```

```
## [1] "1" "2" "3"
```

```
class(x)
```

```
## [1] "character"
```

Objects in R: reading data

- You can import any data file and assign it into an object

```
x <- c(1, 2, 3)
x <- as.character(x)

print(x)
```

```
## [1] "1" "2" "3"
```

```
class(x)
```

```
## [1] "character"
```